

Retrieval-Augmented Generation

1. Introduction

Retrieval-Augmented Generation (RAG) is a technique that improves Large Language Model (LLM) responses by grounding the model in real, trusted information sources. Instead of allowing the LLM to answer solely from its internal parameters, RAG retrieves relevant context from external knowledge bases such as PDFs, databases, websites, or enterprise documents, and supplies them to the model.

This reduces hallucination, improves accuracy, and enables domain-specific responses without retraining the model.

2. Why RAG Is Needed

Pure LLMs rely on:

- Pretrained patterns
- Probabilistic reasoning
- Knowledge up to a cutoff

They cannot access:

- Private company data
- Updated information
- Regulatory documents
- Customer-specific instructions

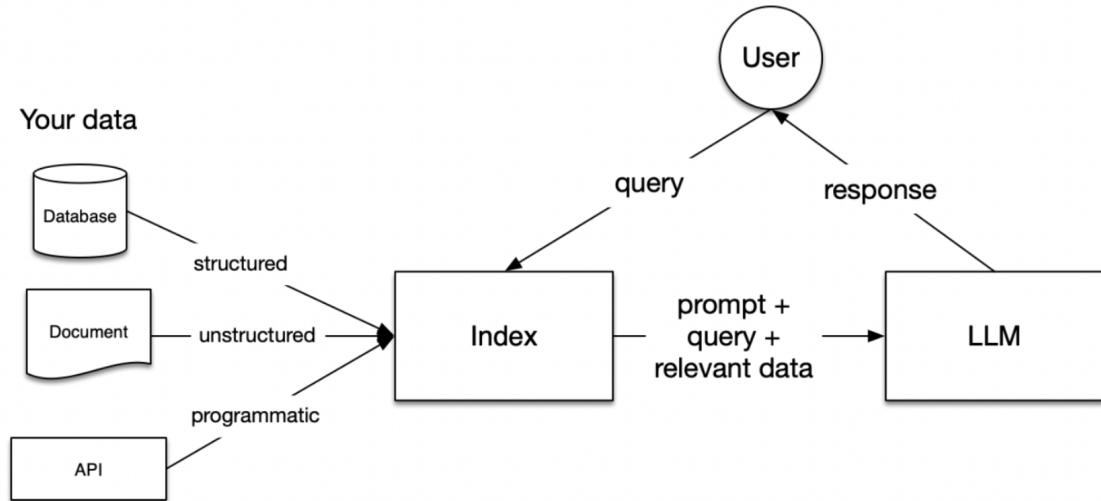
RAG solves this by injecting external knowledge into the prompt at runtime.

3. High-Level Workflow

A RAG pipeline follows this sequence:

1. The user asks a question.
2. The system converts the question into an embedding (a vector representation).
3. The vector is used to search a vector store (FAISS, Chroma, Pinecone, etc.) for semantically similar documents.
4. Relevant documents are retrieved.
5. The LLM receives both the query and the retrieved context.
6. The LLM generates an answer grounded in real information.

This ensures responses are accurate and verifiable.



4. Detailed Mechanics of Each Stage

4.1 Input Question

The user provides a natural-language question. Example:
"What are the steps to configure system access for new employees?"

This becomes the starting point for the retrieval process.

4.2 Embedding Generation

The system transforms the query into an embedding using an embedding model. An embedding is a numerical vector representing semantic meaning.

Important characteristics:

- Similar meanings are close in vector space.
- The embedding model must match the one used for indexing documents.

4.3 Vector Search in the Knowledge Base

All documents stored in your knowledge base are pre-converted into embeddings.

During retrieval:

- The query vector is compared with stored document vectors.
- Techniques like cosine similarity or inner product determine relevance.
- The system selects the top-k most similar chunks.

This step is the core of RAG retrieval.

4.4 Chunking and Metadata

Documents are chunked into smaller units (200–1,000 tokens each).
Each chunk is stored with metadata such as:

- Source file path
- Page number
- Section title

Chunking ensures better recall and precision during search.

4.5 Context Assembly

The retrieved chunks are combined into a prompt block known as the "context window".
The system must:

- Limit context length to avoid token overflow
 - Remove duplicates
 - Maintain chronological or structural order if needed
-

4.6 Prompt Construction

The final prompt given to the LLM includes:

- System instruction
- User query
- Retrieved context

Example structure:

System: "Use only the context below to answer the question."

Context: [All retrieved chunks]

User: [Original question]

This ensures the LLM answers based on real data.

4.7 Generation

The LLM generates an answer grounded in the provided context.

Characteristics of RAG answers:

- Factually anchored
 - Citeable (if required)
 - Less hallucination
 - Context-specific
-

4.8 Optional Enhancements

RAG can be extended with advanced components:

1. Query rewriting
Improves retrieval quality by expanding or reformulating the question.
2. Multi-hop retrieval
Retrieves related context iteratively.
3. Re-ranking
Uses cross-encoders to choose the most relevant chunks.
4. Validation
Checks if the generated answer is grounded in context.

These turn RAG into an enterprise-grade system.

5. RAG vs Model Training

Dimension	RAG	Model Training
Cost	Low	High
Speed to Update	Instant	Slow
Freshness of Knowledge	Real-time	Requires retraining
Hallucination Risk	Low (grounded)	Higher

Dimension	RAG	Model Training
Auditability	High	Low
Data Privacy	Strong isolation	Risk of leakage
Scalability	Unlimited	Model size limited
Control	Explicit	Implicit
Maintenance	Easy	Complex

6.